


Wing Section Studio

User & Reference Guide

Designing multi-element wing sections in ground effect — configuration, analysis, target-downforce optimization, airfoil shape refinement, manufacturing preparation, and CAD-ready export — together with a complete account of how the aerodynamic engine works.

Part I — Using the application · Part II — How it works
Revision July 2026



Contents

1 Overview	3
2 The workflow at a glance	3
3 Getting started	4
4 Step 1 — Configure the section	4
5 Step 2 — Analyze	5
6 Step 3 — Optimize to a target	6
7 Step 4 — Airfoil shape refinement	7
8 Step 5 — Manufacturing preparation	7
9 Support tools — screener and polars	9
10 Step 6 — Export to CAD	9
11 Operating maps and the RANS handoff	9
12 Architecture	11
13 The inviscid panel method	11
14 Ground effect and the correction model	13
15 Viscous section data	13
16 Element loading budget	14
17 The drag model	14
18 The airfoil spec system	15
19 Manufacturing geometry	15
20 Shape refinement mathematics	16
21 The optimizer	16
22 Numerical accuracy and validation	17
23 Model scope and limitations	17
24 Reference	17

PART I

Using the application

1 Overview

Wing Section Studio designs the two-dimensional section of a multi-element racing wing operating in ground effect — the front wing of a Formula-style car being the canonical case. It covers the fast, iterative phase of wing design: choosing airfoils, arranging one to four elements with correct slot geometry, predicting downforce and drag at the operating point, optimizing the arrangement (and, if you wish, the airfoil shapes themselves) to a target load, preparing the section for manufacture, and exporting CAD-ready geometry.

The application is a screening and optimization layer, not a replacement for high-fidelity analysis. Its solver produces exact inviscid solutions of the full multi-element interaction in milliseconds, corrected to realistic values by transparent, user-adjustable calibration factors. Designs shortlisted here are meant to be confirmed with RANS CFD or wind-tunnel measurement, and the calibration factors re-tuned against those results. Part II of this guide explains every step of that computation.

Capabilities at a glance

- One to four elements; 2,174 bundled UIUC airfoils, 4-digit NACA generation, and Selig `.dat` upload.
- Slot geometry set directly as gap and overlap (percent of chord); placement is solved so the achieved values equal the requested ones.
- Coupled inviscid analysis of all elements in free air and in ground effect, with per-element viscous loading limits and a realistic profile-plus-induced drag estimate.
- Target-downforce optimization over stack angle, deflections, slot geometry, chords, airfoil choice, and airfoil shape — returning a set of distinct candidate designs.
- Manufacturing preparation: buildable trailing edges and thickness checks, folded through the whole pipeline so what you analyze is what you cut.
- Operating maps sweeping ride height or speed around the design point, with the current point and the downforce peak marked.
- Export to DXF, Selig `.dat`, XYZ points, CSV, SVG, and a JSON manifest, in both the installed and upright orientations — plus a ready-to-run OpenFOAM 2D RANS case.

2 The workflow at a glance

Analysis pipeline (`analysis.analyze`)



Figure 1 — Inside a single analysis: validated inputs become a placed, solved stack, an inviscid pair of solutions, a corrected estimate, and viscous loading and drag — described in full in Part II.

A typical session runs: pick a starting template; set the operating point (speed, ride height, chord, span); adjust elements and slots while watching the live drawing and loading meters; press **Analyze** to get forces; optimize to a target; optionally refine the airfoil shapes; enable manufacturing preparation; and export. Analysis takes well under a second, so the intended rhythm is conversational — adjust, analyze, repeat.

3 Getting started

Launch by double-clicking **Wing Section Studio** (desktop shortcut or project folder). A window opens immediately with a loading screen while the aerodynamic models start — a few seconds — then the workspace appears. Everything runs locally in one process; closing the window shuts it down. The last session's configuration is saved automatically and restored on the next launch.

Configuration	Operating point, target downforce, the elements, the manufacturing group, and an advanced group of air properties and model-calibration factors.
Elements	One card per element: airfoil, chord, deflection, slot gap and overlap, with live Reynolds number and as-built badges.
Section drawing	A dimensioned technical drawing in the as-driven orientation, with ground hatching, ride-height and length dimensions, slot callouts, and the centre-of-pressure mark. Scroll to zoom, drag to pan.
Results	Estimated downforce against target, coefficient and drag tiles, per-element loading meters, and plain-language warnings.
Detail tabs	Pressure distributions, section polars, the optimizer, the airfoil screener, and export.

Two orientations, one geometry

The drawing shows the wing the way it runs on the car — inverted, above the road — because that is the orientation the analysis uses. The **Upright** toggle mirrors it into the convention airfoil catalogs and CAD sketches use. Both are exact rigid transforms of the same section; nothing about the physics changes.

4 Step 1 — Configure the section

Operating point

Speed	Design airspeed. Forces scale with dynamic pressure; each element's Reynolds number follows from speed and its chord.
Ride height	Clearance from the road to the lowest point of the section, in millimeters — the dominant ground-effect parameter.
Main chord / Span	Main-element chord and wing span. Span enters the reference area and the induced-drag aspect ratio.
Stack angle	Incidence of the whole section. Positive is nose-up — more downforce — the same sense as flap deflection.
Ncrit	Transition criterion for the viscous model. 9 is a clean tunnel; 4–7 represents on-track turbulence.
Air properties & calibration	Density, viscosity, the two calibration factors of the estimate — the free-air viscous realization and the ground-gain realization factor, left blank for the automatic ride-height curve (section 14) — the finite-span lift efficiency, the span efficiency for induced drag, and panel resolution.

Elements and slot geometry

Each element card selects an airfoil — type to search the library, press **.dat** to load a coordinate file — and, for flaps, sets chord (percent of main chord), deflection (degrees, trailing edge down), and the two slot parameters used throughout the high-lift literature:

Slot geometry — gap (nearest clearance) and overlap (horizontal tuck)




Figure 2 — Slot definitions on a real placement. Gap is the smallest clearance from the previous trailing edge to the flap surface; overlap is the horizontal tuck of the flap nose under that trailing edge.

Placement is solved so the section achieves exactly the requested gap and overlap. Healthy starting ranges are 1–2.5 %c gap and 2–4 %c overlap; the card badges always show the achieved values, and the analysis warns when a slot chokes (gap below 0.5 %c) or a requested gap cannot be reached. The drawing updates live with every edit.

5 Step 2 — Analyze

Analyze section solves the coupled inviscid flow around all elements twice — once in free air and once in ground effect — and combines the results with per-element viscous data into the numbers on the right-hand panel:

- **Estimated downforce**, compared against the target as an on-target / over / under pill. The small information button beside the label opens the exact computation.
- **Coefficient tiles** — the corrected estimate alongside the raw inviscid coefficients (ground and free air) it derives from, and the inviscid ground-effect gain.
- **Drag estimate** — induced plus profile, with the split shown beneath the total and the induced-drag details on hover.
- **Element loading meters** — each element's working lift against the isolated maximum of its airfoil at its own Reynolds number. Warnings begin at 90 % of the limit, criticals above 110 %: the classical budget for slotted high-lift systems. A second check watches the realized ground-effect operating point (section 16).
- **Pressure tab** — the inviscid surface-pressure distribution of every element, the physics-level view of how the elements share load.




Figure 3 — Loading meters from a real analysis. The main element runs above its isolated maximum (critical, red) — normal for a driven main plane — while the flap sits comfortably under (green). The dotted line is the 90 % warning threshold.

6 Step 3 — Optimize to a target

Set the target downforce and choose the free variables: stack angle, flap deflections, slot gap and overlap, flap chords, the airfoil of each element, and the airfoil shape (section 7). The search drives predicted downforce to the target while penalizing drag, slot geometry outside the healthy band, element overload, and intersecting geometry.




Figure 4 — A real optimization run: downforce converging to the 250 N target (left) and the objective descending on a log scale (right).

Candidates, not a single answer

A finished run returns the best design plus up to three **candidate designs** — distinct set-ups that also hit the target, chosen to be genuinely different from one another (different slot geometry, incidence, or airfoils) rather than trivial variations around one optimum. Each candidate card shows its own analyzed downforce, drag, and lift-to-drag, and applies to the configuration with one click — so build-tolerance and packaging trade-offs stay visible instead of being hidden inside a single number.

Effort and reproducibility

Fast / Standard

Stop as soon as the target is reached cleanly — a few seconds to half a minute. Quick, but because the search halts in the first basin that meets the target, two runs from slightly different starting points can land on different (similarly-good) designs.

Thorough

Searches the whole space before refining, with a wider population and a multi-start final polish at full solver resolution. It takes minutes, and repeated runs converge to the same lowest-drag design — use it for a design you intend to keep.

When a target is unrealistic in either direction, the result panel says so in plain language — for example, that the stack met the target only by parking its loading variables at their minimum, so the target could be raised or an element removed. **Apply to configuration** writes the chosen design, including any airfoil or shape changes, back into the element cards and re-analyzes it.

7 Step 4 — Airfoil shape refinement

Airfoil selection chooses the best existing section for each element. **Airfoil shape** goes further: it lets the optimizer modify that section — pushing camber where the loading wants it, thinning or thickening the envelope — to squeeze out more performance. Enable it as a free variable alongside the others.




Figure 5 — Shape refinement. Three camber levers (front, mid, aft loading) plus a thickness scale (left) modify whatever section the element uses (right); the modification is applied to a real S1223 here.

Each element gets three camber adjustments (near 25, 55, and 80 % chord — the classical front-, mid-, and aft-loading levers) and one thickness scale. Every constraint still holds: manufacturing trailing-edge preparation wraps the modified section, the buildable-thickness floor limits how far it may be thinned, and the stall budget uses the modified section's own polar. Shaped sections carry through polars, exports, and saved projects; re-running the optimizer keeps refining the current shape rather than stacking changes.

Shape refinement is slower

Every new shape needs fresh aerodynamic evaluation (there is no library polar to reuse), so shape runs are somewhat slower per evaluation — roughly 1.3–1.7× — than geometry-only runs. Pair it with the **Thorough** effort setting when you want a final, reproducible shape, and expect a run to take a few minutes.

8 Step 5 — Manufacturing preparation

Theoretical airfoil coordinates close to a knife edge at the trailing edge — zero thickness. No layup, print, or hot-wire cut can produce that. Enable **Manufacturing** to open every element's trailing edge to a minimum buildable thickness in millimeters, and to check that each section is thick enough to make.

Trailing-edge treatments — a knife edge cannot be built




Figure 6 — The two trailing-edge treatments applied to a real section, zoomed on the aft. Thicken adds material over the rear while keeping the chord; Cut off truncates the tip and rescales.

Thicken	Adds thickness over the rear of the section, keeping chord and camber. Costs a fraction of a percent of downforce at a typical 1–1.5 mm trailing edge (about 0.3 % on the baseline section), rising to roughly 2 % at 3 mm. This is the default and the recommended aerodynamic preparation.
Cut off	Truncates the section where it reaches the target thickness and rescales to the original chord — for when a part or mold must literally be cut back. Thin-tailed sections lose far more downforce than the thicken treatment — roughly 17 % on the baseline section at 1.5 mm — and the tooltip says so.
Min thickness	An optional buildable-minimum check on the whole section. It warns when an element's thickest point is below the minimum and keeps thinner airfoils out of the optimizer and screener — it never silently reshapes a section you chose.

Thicken does more than open the trailing edge: it also raises the thickness of any station behind the section's nose that would otherwise be thinner than the trailing edge itself — a thin sail-like section would otherwise keep an unbuildable waist ahead of its blunt tail. A section whose thickest point is not meaningfully above the requested trailing edge is flagged **critically as a plate**: the treatment still opens and floors it, but the result is a near-constant-thickness plate the tool warns you about rather than a section you would have chosen (Figure 7).




Figure 7 — As-built thickness distributions. Left: a buildable section, floored so nothing behind the nose is thinner than the 3 mm trailing edge. Right: a section thinner than the requested edge everywhere — a plate: still treated, but flagged critically.

Manufacturing preparation is fully reversible

Turning the checkbox on never alters your configuration — the treated shapes are derived on the fly, and turning it off returns the exact original geometry. Each element card shows the as-built trailing-edge thickness, maximum thickness, and the thinnest station behind the nose (the **waist** badge). Every downstream step — the solver, polars, the optimizer, and all exports — uses the as-built shapes, so what you cut is what was analyzed. If a design is destined for manufacture, keep the box on while optimizing so the search tunes the real part.

9 Support tools — screener and polars

Airfoil screener

The screener sweeps all bundled sections (plus any uploads) at a chosen Reynolds number and tabulates maximum lift, best efficiency, drag at a reference lift coefficient, thickness, camber, and the surrogate model's confidence. Columns sort on click; **Use** assigns a section to an element. Screening the full library takes a few seconds and is cached, so re-sorting and re-screening are instant. When manufacturing preparation is on, the thickness filter is pre-set so sections too thin to build are excluded up front.

Polars

The polar tab plots each element's lift curve and drag polar from the neural-network surrogate (trained on XFOIL data; instant), for the section as actually built. An **Add XFOIL reference** button runs the locally installed XFOIL executable, when present, for an independent cross-check — the two typically agree within a few percent below stall.

10 Step 6 — Export to CAD

DXF	True-scale millimeters, one closed contour per element on its own layer, as smooth splines (for lofting) or exact polylines. Opens directly as a sketch in SolidWorks, Fusion, Onshape, NX.
Complete bundle (ZIP)	DXF in both orientations, per-element XYZ point files, Selig .dat contours (main-chord units, for meshing and panel workflows), CSV, an SVG drawing, and a JSON manifest with the full configuration and analysis snapshot.
SVG drawing	True-scale vector drawing with the ground line, for documentation or 1:1 print checks.
Point table (CSV)	Every contour point in millimeters with element metadata.
OpenFOAM case	A complete, ready-to-run 2D RANS case — mesh, boundary conditions, solver settings, and a run script — at a chosen mesh resolution. The handoff workflow is section 11.

Exports can be written in the as-driven orientation (ground at $y = 0$, ready for CFD domains) or upright (the CAD sketch convention). Each export is saved to the project's **exports/** folder with a timestamped name, and a dialog shows the exact location with a **Show in folder** button; a browser download of a copy is offered alongside. With manufacturing preparation enabled, every exported contour is the as-built profile, and the manifest records the achieved trailing-edge thickness per element.

11 Operating maps and the RANS handoff

Two closing tools round out the workflow. The operating maps show how a finished design behaves away from its design point — a wing tuned at one ride height rides at many — and the OpenFOAM export hands a shortlisted section to a genuine RANS solver for confirmation.

Operating maps

The **Maps** tab (beside the Airfoil screener) sweeps the current design across ride height or speed — by default 10 to 120 mm of ride height in 23 steps, up to 40 points — and plots downforce and lift-to-drag against the

swept value, marking the current operating point and the downforce peak. A ride-height sweep re-solves the flow at every point and takes a second or two at the sweep's capped panel resolution (reported under the charts; **Analyze** always runs at full resolution). A speed sweep is near-instant: the inviscid solution does not depend on speed, so only the Reynolds numbers and the dynamic pressure vary.




Figure 8 — The Maps tab's ride-height sweep of the default template, generated by the same sweep the application runs. Downforce peaks near 20 mm and falls closer to the road; the dotted line is the configured operating point.

The force-reduction peak the correction model predicts (section 14) is now directly visible instead of being a property taken on trust: on the default two-element template the peak sits near 20 mm ride height — running lower loses downforce. Swept points that push an element past twice its isolated stall limit are flagged beneath the charts: the model is optimistic there, so treat those downforce values as upper bounds.

The OpenFOAM handoff

Everything this tool reports screens and ranks; RANS decides. The **OpenFOAM case** card in the Export tab writes a complete, ready-to-run two-dimensional RANS case for the current section: a mesh with boundary layers resolved down to the wall on the wing surfaces and on the ground beneath the wing, the steady $k-\omega$ SST setup, a ground plane moving at the freestream speed (the road, in the wing's frame), and automatic force-coefficient extraction. The reported lift coefficient is downforce-positive and referenced to the main chord — directly comparable to the coefficients this guide has used throughout.

Coarse (~15k cells)	A first look — does the flow stay attached, is the case healthy. Runs in minutes.
Medium (~40k cells)	The default; the standard confirmation run for a shortlisted design.
Fine (~90k cells)	Final numbers, and the grid-sensitivity check for a design you intend to commit to.

The workflow: generate the case (it lands in the exports folder like any other export), then run it from the case folder under WSL — the one-time OpenFOAM installation is described in the repository README:

```
ws1 -d Ubuntu -- bash run.sh
```

The script prepares the OpenFOAM environment, converts and checks the mesh, runs the solver, and writes the final coefficients to `results.txt`; each case's `README.txt` records its exact numbers and conventions. Compare the RANS downforce with the estimate here, and re-tune the calibration factors (section 14) against it — closing exactly the loop those factors were built for.

PART II

How it works

12 Architecture

The application is one local process. A thin desktop shell (a WebView2 window) hosts a browser user interface built from plain HTML, CSS, and JavaScript with no build step. That interface talks to an in-process web server over a loopback-only HTTP connection; the server exposes every capability as a small REST API and delegates all real work to a framework-free Python core.

Architecture — one local process, aerodynamics isolated in app.core




Figure 9 — The whole application in one process. All aerodynamics lives in `app.core`, framework-free Python, reachable identically from the window or from scripts.

Everything aerodynamic lives in `app.core`: the panel solver is pure NumPy, with NeuralFoil and AeroSandbox for viscous data and geometry, SciPy for the optimizer's global search, and Matplotlib and ezdxf for geometry tests and DXF output. No web framework reaches into it. The web layer is a wrapper: the same functions that draw the live section and run the optimizer are importable directly for scripting, and the regression suites exercise them without a server. The server binds only to the loopback address and rejects any request whose Host header is not local, closing the one route by which a web page could otherwise reach it.

Frames of reference

Two coordinate frames recur throughout. The **design frame** is upright — lift positive-up, the main leading edge at the origin, the main chord equal to one. The **installed frame** is the design frame flipped so downforce points down, with the ground plane at $y = 0$ and the whole section translated so its lowest point sits at the ride height. The panel solver runs in the installed frame; the two orientations you can draw and export are exactly these frames. All stack coefficients are referenced to the main chord; per-element lift coefficients are referenced to each element's own chord.

13 The inviscid panel method

The core solver is a multi-element Hess–Smith panel method. Each element contour, in Selig ordering, is discretized into flat panels. Every panel carries a constant-strength source distribution (one unknown per

panel), and all panels of a given element share a single vortex strength (one more unknown per element). Sources set the shape of the flow; the per-element circulation sets its lift.

Flat-panel discretization: nodes, panels, and outward normals at collocation points




Figure 10 — A real section discretized into flat panels. The solver enforces the boundary conditions at the collocation point of each panel, using the outward normals shown.

The unknowns are fixed by two sets of conditions. At the collocation point of every panel, **flow tangency** requires zero velocity normal to the surface — the solid-wall condition. For each element, one **Kutta condition** requires the flow to leave the trailing edge smoothly, expressed as equal-magnitude tangential velocities on the two panels meeting there. Together these give a single dense linear system spanning all elements at once, so every slot and stack interaction is captured exactly — there is no assumption that the elements are independent.

Ground effect by images

The ground plane is introduced by the method of images: the velocity the real system induces at a point equals what a mirror system induces at the mirrored point, under a single reflection identity. Reflecting a source keeps its sign; reflecting a vortex flips it; both facts fall out of the same identity, so the ground adds no new unknowns — only an extra influence term. This is important because the naive alternative (a lift-from-circulation shortcut) is simply wrong near the ground, where image-induced velocities are large.

Forces from pressure

Once the strengths are known, the surface pressure coefficient at each panel follows from its tangential velocity, and the forces and moment come from **integrating that pressure** over the surface — not from the circulation. Pressure integration remains valid in ground effect, where the circulation-based shortcut over-predicts downforce by roughly a factor of two because it ignores the image-induced velocity field. The moment is reported nose-up-positive, and the centre of pressure follows from it. A small numerical drag residual (exactly zero for ideal inviscid flow) is reported as a discretization diagnostic.




Figure 11 — Real inviscid surface pressure for a two-element stack, free air (dashed) versus ground effect (solid). Ground effect deepens the suction on the main element's lower surface — the venturi that makes downforce.

14 Ground effect and the correction model

The panel solution is exact for inviscid flow, but inviscid ground effect is unphysical in the limit: as the ride height shrinks, the predicted venturi gain grows without bound, while a real flow chokes on boundary-layer growth in the gap. The tool therefore reports a corrected estimate built from transparent, user-adjustable factors:

```
C_est = eta_visc · [ C_free + G_real ]
G_real = cap · tanh( k_g · (C_ground - C_free) / cap ) · tanh( (h/c) / 0.045 )
cap = 3.0 · |C_free|
```

Here C_{free} and C_{ground} are the exact inviscid downforce coefficients in free air and in ground effect. η_{visc} (default 0.85) is the free-air viscous realization — how much of the inviscid free-air load a real viscous flow achieves. k_g is the fraction of the additional ground-effect gain that is realized; at ordinary ride heights both saturation terms are near-linear and near-one, so the model reduces to the familiar $\eta_{\text{visc}} \cdot [C_{\text{free}} + k_g \cdot (C_{\text{ground}} - C_{\text{free}})]$. The first saturation caps the realized gain at three times the free-air load; the second chokes it off below $h/c \approx 0.045$, where the real venturi flow stalls on boundary-layer growth. Left blank, k_g follows an automatic ride-height curve:

```
k_g = 0.85 · tanh( (1.4 / 0.85) · h/c )
```




Figure 12 — Left: the automatic realization-factor curve, which preserves a calibrated small-height slope and an 0.85 ceiling but vanishes at the ground. Right: on a real stack, the raw inviscid gain diverges as ride height shrinks while the realized estimate peaks and then falls toward the road.

Together the saturations keep the calibrated slope at racing ride heights while genuinely reversing at the bottom of the travel: the estimate peaks near $h/c \approx 0.065$ on the two-element baseline and decreases as the wing drops further — the force-reduction behaviour measured on wings in strong ground effect — instead of growing without bound. Both knobs are exposed in the calibration group beside the raw inviscid numbers they modify: η_{visc} directly, and k_g , which follows the automatic curve when left blank and is pinned when a value is entered — useful for matching RANS or tunnel data at a known ride height. The analysis always reports which source it used. The cap ratio and choke height (gain_cap_ratio , choke_h_c) are configuration fields with conservative defaults, adjustable at the config/API level.

15 Viscous section data

The panel method knows nothing of viscosity, so the section-level viscous quantities — maximum lift, profile drag, the angle for a given lift — come from NeuralFoil, a neural-network surrogate trained on large sweeps of XFOIL solutions. It evaluates a complete polar in milliseconds, which is what makes library-wide screening and thousands of optimizer evaluations feasible, and it reports a **confidence** value that falls for unusual sections

and operating points.

A locally installed XFOIL executable — dropped into the application's `xfoil/` folder; it is not redistributed with the tool — serves as an optional independent cross-check through the same interface; it is slower and single-element, but it is the reference the surrogate approximates. Polars are cached by section, Reynolds number (bucketed to three significant figures), transition criterion, and model size, so a repeated operating point is free. The optimizer and the final analysis read the same fast surrogate; the polar tab reads a larger, more accurate one. The optimizer's speed comes instead from coarser panel resolution during the search.

Where to be sceptical

The surrogate is most trustworthy for conventional sections at moderate Reynolds numbers. Very unusual shapes or very low Reynolds numbers deserve the XFOIL cross-check. When a section's polar never stalls within the analyzed angle range, the reported maximum lift is a lower bound rather than a true stall value, and the loading warnings say so.

16 Element loading budget

A slotted high-lift system can carry more total lift than the sum of its isolated elements, but each element still has a viscous ceiling. Following classical high-lift practice, the tool budgets each element's **free-air** inviscid load — knocked down by the viscous realization factor — against that element's isolated maximum lift coefficient, evaluated at its own Reynolds number (Figure 3). Warnings begin at 90 % of the limit and criticals at 110 %, since slotted elements tolerate somewhat more than isolated maximum.

Free-air load is used for that budget deliberately: raw inviscid ground-effect loads grow without bound near the road and would make the check meaningless. A second indicator watches each element's **realized ground-effect operating point** — the same lift coefficient the profile-drag lookup uses, so it is consistent with the reported estimate by construction. Sections near the ground have been measured carrying roughly up to twice their isolated maximum before the flow gives up, so an element past that allowance draws a screening-validity warning: treat the downforce estimate as optimistic there, and the profile drag — capped at the pre-stall polar — as understated. Each element's inviscid ground-to-free multiplier is reported alongside, so the loading picture stays honest while the ground gain remains visible.

17 The drag model

Reported drag has two parts. **Profile drag** is each element's own section drag, read from its polar at the lift coefficient it actually works at in ground effect (its free-air load plus the realized share of its ground-effect gain, viscous-realized), then summed over the elements weighted by chord. **Induced drag** is computed from the very downforce the tool reports, so lift and induced drag are always consistent:

$$CD_i = C_L^2 / (\pi \cdot AR \cdot e) \cdot \phi_i$$

where `AR` is the wing aspect ratio, `e` the span-efficiency input (endplates push it toward one), and `phi` the McCormick ground-effect factor evaluated at the section's mid-height — wings near the ground shed much weaker trailing vorticity, so induced drag falls as the road approaches:

$$\phi_i = (16 h/b)^2 / (1 + (16 h/b)^2)$$





Figure 13 — The McCormick factor. Close to the ground, the trailing vortex system weakens and induced drag drops sharply.

For a loaded finite-span front wing, induced drag dominates the total — often by an order of magnitude over profile drag — so leaving it out (as a profile-only figure would) produces fantasy efficiency numbers. Interference and mounting drag are not modeled.

18 The airfoil spec system

Every airfoil in the application — a library section, an upload, a manufactured section, a shape-modified section — is named by a single **spec string**, and one resolver turns any spec into coordinates. Modifiers are prefixes that wrap a base spec, so they compose:

Airfoil spec resolution — one string names any section, raw or modified



Wrappers nest outward: a shaped, manufactured library section is shape:....mfg:....s1223 — every layer is a first-class airfoil to the solver, polars, screener and exporters.

Figure 14 — Spec resolution. Each modifier unwraps to its base and re-applies its transformation, so a shaped, manufactured library section is one string that every part of the system understands.

Because resolution is centralized, a modified section is a first-class airfoil everywhere with no special cases: the panel solver, the NeuralFoil polars, the XFOIL cross-check, the screener, and every exporter all see the same as-built coordinates. Modifier parameters are quantized (trailing-edge gaps and shape parameters each snapped to a fixed grid) so the spec strings — and the repanel and polar caches keyed on them — stay stable while the optimizer sweeps continuous values.

19 Manufacturing geometry

The **thicken** treatment adds thickness symmetrically about the camber line, blended in over the rear of the section with a smooth ramp, so the added opening between the two surfaces at the trailing edge equals the requested gap exactly. It then applies a **thickness floor**: from the first station where the section reaches the requested gap, back to the trailing edge, no station is left thinner than the gap. This is what removes the mid-chord waist that a naive ramp leaves on thin sections (Figure 7). The treatment works on a densely repaneled copy of the base section, and the floor is guaranteed on the geometry every consumer actually sees: the floored contour is verified through a spline-repanel round-trip, with the floor raised adaptively until it survives the fit.

The **cut off** treatment truncates the section at the station where it reaches the target thickness — solved so that after rescaling to the original chord the trailing edge is exactly the requested thickness — and closes it with a flat base. It applies no floor, because its purpose is to cut a real part back; instead the report measures any remaining waist and warns. Because camber is preserved by the symmetric construction, the treatment is a **manufacturing operation, not a hidden aerodynamic re-trim**: shifting material onto one surface instead would just be a camber change, which the optimizer's real variables already control.

Exports close the blunt base explicitly. In the spline DXF the base is a separate straight line between the two surface endpoints, keeping the corners sharp instead of rounding them; the polyline DXF and the point files carry the exact open contour. The manifest records the achieved trailing-edge thickness, maximum thickness, and the measured minimum thickness behind the nose for every element.

20 Shape refinement mathematics

Shape refinement modifies a section with a small, smooth, airfoil-preserving parameterization. The camber line is displaced by a sum of three **Hicks–Henne bump functions** centred near 25, 55, and 80 % chord — each a smooth hump that is zero at both ends of the chord and peaks at its centre — and the thickness envelope (each surface's offset from the camber line) is scaled by a single factor. Four numbers per element (Figure 5).

This choice is deliberate. Four smooth parameters keep the added search dimensions cheap and keep every generated shape recognizably an airfoil, so the NeuralFoil surrogate stays inside the space it was trained on and its confidence remains meaningful — a free-form coordinate parameterization would wander into shapes the surrogate cannot rate. The shape wraps as a spec modifier (section 18), so manufacturing preparation applies on top of the shaped section, and the buildable-thickness minimum limits how far the thickness scale may thin an element. Re-optimizing a shaped design unwraps and seeds from its current parameters rather than nesting a second modification.

21 The optimizer

The optimizer treats the design as a vector of continuous variables and minimizes a single objective. Continuous quantities — stack angle, deflections, slot gap and overlap, chords, shape parameters — map directly. **Airfoil choice** is encoded as one continuous variable per element mapped onto a shortlist of candidate sections ranked by maximum lift, so that neighbouring values are similar airfoils and the search stays meaningful. The shortlist is built by the screener at each element's Reynolds number, includes the currently selected section (unless the candidate pool is restricted to your uploads), and — when manufacturing sets a buildable minimum — excludes sections too thin to make at that element's chord.

Objective and constraints

The objective is dominated by the squared relative error to the target downforce, so the search first and foremost hits the number you asked for. Added to it are a weighted drag term (the adjustable drag penalty, trading residual drag against exactness on target) and soft penalties that keep candidates inside the healthy high-lift envelope: slot gap and overlap bands, element loading below about 105 % of each airfoil's isolated maximum, and a hard penalty on intersecting geometry or a failed solve. The loading penalty is weighted so that an honest miss of the target beats a design that only reaches it by driving an element deep into stall.

Search, candidates, and reproducibility

The search is a seeded global phase (differential evolution) followed by a local refinement (Nelder–Mead), evaluating a few hundred to a few thousand complete analyses. Identical inputs give identical results — the search is fully seeded. Every feasible evaluation is archived; at the end, the best design plus a few maximally-different on-target designs are selected from that archive as the candidate set, each re-analyzed at full resolution.

Fast and Standard runs stop as soon as the target is met, which is quick but means the design depends on which basin the search reached first. Thorough runs remove that dependence: the global phase runs to completion with a wider population, and the final polish runs from several diverse starting points at full panel resolution, with the reported best required to come from full-resolution evaluations. The result is that two Thorough runs from different starting configurations converge to the same lowest-drag design, where Fast runs might differ.

22 Numerical accuracy and validation

The panel method is validated against an independent linear-vorticity formulation (with forces taken by pressure integration) and against XFOIL's inviscid solution on single elements, agreeing to within a few percent in free air and in ground effect. The image system is checked against an explicitly modelled mirror geometry, and the slot gap/overlap solver against direct geometric measurement of the placed sections.

The repository ships regression suites that exercise the solver physics, the manufacturing geometry (including the thickness floor and plate detection), the optimizer (including determinism and Thorough reproducibility), the shape system, and the API's failure behaviour under hostile input. They run from a single command and form the contract the application is held to; every figure in Part II of this guide is generated from the same core the suites test.

What the numbers are for

Treat every downforce and drag figure as a screening and ranking value. The tool is calibrated to land a typical two-element section at racing ride heights in the right physical band, but it resolves neither the boundary layer, slot merging and separation, nor three-dimensional effects. Confirm a shortlisted design with RANS CFD or measurement, and re-tune the calibration factors against those results, before committing to it.

23 Model scope and limitations

- **Two-dimensional section physics** with classical finite-span corrections. Spanwise design — endplates, footplates, local twist — is out of scope.
- **Inviscid pressure field.** Boundary-layer displacement, slot merging, and separation are represented only through the calibrated corrections and the loading limits, not resolved.
- **Calibrated ground effect.** The correction factors are tuned to a physical band, not derived; strong ground effect in particular should be confirmed against higher-fidelity data.
- **Surrogate viscous data.** Section polars come from a neural surrogate of XFOIL; its confidence is reported and a cross-check is one click away, but unusual sections at low Reynolds numbers deserve scrutiny.
- **No interference or mounting drag**, and no compressibility (irrelevant at these speeds).

24 Reference

Project files

Save and Open write a portable JSON project containing the full configuration, target, and any uploaded airfoil geometry (including inside shape and manufacturing wrappers). The working session is also autosaved server-side, in `app_data/` beside the application, and restored on launch.

Exports folder	Exports are written to exports/ beside the application, with timestamped names; the export dialog links to the exact file.
Automation / API	Every function is available over the local REST API; interactive documentation is served at /api/docs while the application or the headless server is running.
Headless use	The same server runs without a window for scripting: <code>python -m uvicorn app.server:app --port 8642</code> .
Verification	The regression suites run from a single command: <code>python app/tests/run_all.py</code> (it starts its own scratch server for the API checks).
Design record	Every design decision behind the application is recorded, with the options considered, in DECISIONS.md at the repository root.

Wing Section Studio — User & Reference Guide. Every figure in this document is generated directly from the application's aerodynamic core, so the guide tracks the shipping behaviour.